

# 6G\_IRS\_Keras

May 25, 2022

```
[1]: !pip install -q cleverhans
!pip install -q plot_keras_history
!pip install -q keras-tuner
```

```
[2]: import scipy.io
import tensorflow as tf

from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
from scipy.linalg import norm

from tqdm.keras import TqdmCallback
from tqdm.notebook import tqdm
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split

from plot_keras_history import plot_history
import matplotlib.pyplot as plt
import keras_tuner as kt
from keras.models import load_model

from sklearn.metrics import mean_squared_error
from cleverhans.tf2.attacks.projected_gradient_descent import ↵
    ↪projected_gradient_descent, fast_gradient_method
from cleverhans.tf2.attacks.basic_iterative_method import basic_iterative_method
from cleverhans.tf2.attacks.carlini_wagner_l2 import carlini_wagner_l2
from cleverhans.tf2.attacks.momentum_iterative_method import ↵
    ↪momentum_iterative_method
from tqdm.keras import TqdmCallback
import numpy as np

import seaborn as sns
import pandas as pd

from keras.models import Sequential
from keras.layers import Dense
```

```

HYPARAMETER_TUNING = False
TRAIN_MODEL = False
GOOGLE_COLAB = True

PROJECT_FOLDER = ''

```

```

[3]: if GOOGLE_COLAB:
    from google.colab import drive
    drive.mount("/content/gdrive")
    PROJECT_FOLDER = '/content/gdrive/My Drive/ds/IRS/'

```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

```

[4]: X = pd.read_csv(PROJECT_FOLDER + 'xtrain_32_10000.csv', header=None)
    X = X.values

    Y = pd.read_csv(PROJECT_FOLDER + 'ytrain_32_10000.csv', header=None)
    Y = Y.values

    M = X.shape[1]

```

```

[5]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33,
    ↪random_state=42)

```

```

[6]: def get_model(mult_factor):
    model = Sequential()
    model.add(Dense(int(M*mult_factor), input_dim=M, activation='relu'))
    model.add(Dense(int(2*M*mult_factor), activation='relu'))
    model.add(Dense(int(4*M*mult_factor), activation='relu'))
    model.add(Dense(int(4*M*mult_factor), activation='relu'))
    model.add(Dense(Y.shape[1], activation='relu'))

    model.compile(loss='mse', optimizer='sgd', metrics=['mse'])

    return model

```

```

[7]: model = get_model(1.0)
    se = tf.keras.callbacks.EarlyStopping(monitor='val_mse',patience=150,verbose=1,
    ↪restore_best_weights=True)


    history = model.fit(X_train, y_train, epochs=500, verbose=0,validation_split=0.
    ↪2,
    ↪callbacks=[se,TqdmCallback(verbose=1)])

```

0epoch [00:00, ?epoch/s]

0batch [00:00, ?batch/s]

```
[8]: plot_history(history)
plt.show()
```



```
[9]: model.save(PROJECT_FOLDER + 'model.hdf5')
```

## 1 Convert the model Cleverhans compatible model

```
[10]: model_copy = tf.keras.models.clone_model(model)
for layer in model_copy.layers:
    layer.trainable = False

model_copy.add(Dense(2, activation='softmax'))
model_copy.compile(optimizer='adam', loss='categorical_crossentropy',
    metrics=['accuracy'])

y_dummy = np.zeros((X_train.shape[0],2))
mid_range = np.int(X_train.shape[0] / 2.0)
y_dummy[0:mid_range,0] = 1.0

history = model_copy.fit(X_train, y_dummy, epochs=2, verbose=0,
    callbacks=[TqdmCallback(verbose=1)], batch_size=256,
    validation_split=0.33, shuffle=True)
plot_history(history)
plt.show()
```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:9:  
 DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To  
 silence this warning, use `int` by itself. Doing this will not modify any

behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information.


Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
if __name__ == '__main__':
```

```
0epoch [00:00, ?epoch/s]
```

```
0batch [00:00, ?batch/s]
```



## 1.1 Attack the model

```
[11]: from itertools import product
      from random import shuffle

      logits_model = tf.keras.Model(model_copy.input, model_copy.output)
      attacks = ['FGSM', 'BIM', 'MIM', 'PGD', 'CW']
      attacks = ['FGSM', 'BIM', 'MIM', 'PGD']

      eps_vals = [0.01, 0.1, 0.3, 0.5, 0.7, 0.8]

      params_list = list(product(eps_vals, attacks))
      shuffle(params_list)

      mal_diff_distance_list = []
      real_mse_list = []
      mal_mse_list = []
      mal_predicted_diff_list = []
```

```

attack_name_list = []
eps_val_list = []

#for i in tqdm(range(X_test.shape[0])):
for _ in tqdm(range(5)):
    i = np.random.randint(0,X_test.shape[0])
    test_input = X_test[i:i+1,:].astype(np.float32)

    for eps_val, attack in params_list:
        if attack == 'FGSM':
            mal_input = fast_gradient_method(logits_model, test_input,
↪eps=eps_val/20.0, norm=np.inf,
                                                    clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                                    targeted=False)

            mal_input = mal_input.numpy()
        elif attack == 'BIM':
            mal_input = basic_iterative_method(logits_model, test_input,
↪eps=eps_val, norm=np.inf,
                                                    #clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                                    eps_iter=0.01,nb_iter=2000,
                                                    targeted=False)

            mal_input = mal_input.numpy()
        elif attack == 'MIM':
            mal_input = momentum_iterative_method(logits_model, test_input,
↪eps=eps_val, norm=np.inf,
                                                    clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                                    eps_iter=0.01,nb_iter=2000,
                                                    targeted=False)

            mal_input = mal_input.numpy()
        elif attack == 'PGD':
            mal_input = projected_gradient_descent(logits_model, test_input,
↪eps=eps_val, norm=np.inf,
                                                    clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                                    eps_iter=0.01,nb_iter=2000,
                                                    targeted=False)

            mal_input = mal_input.numpy()
        elif attack == 'CW':
            mal_input = carlini_wagner_l2(logits_model, test_input,
                                                    targeted=False, #y=y_output,
                                                    batch_size=512, confidence=10.0,
                                                    abort_early=False, max_iterations=1000,

```

```

clip_min=test_input.min(),clip_max=test_input.
↪max())

mal_diff = mal_input - test_input
mal_diff_distance = norm(mal_diff, ord = np.inf)
mal_diff_distance_list.append(mal_diff_distance)

real_output = y_test[i:i+1,:]
test_output = model.predict(test_input)
real_mse = mean_squared_error(real_output,test_output)
real_mse_list.append(real_mse)

malicious_test_output = model.predict(mal_input)
mal_mse = mean_squared_error(real_output,malicious_test_output)
mal_mse_list.append(mal_mse)

mal_predicted_diff = malicious_test_output - test_output
mal_predicted_diff = norm(mal_predicted_diff, ord = np.inf)
mal_predicted_diff_list.append(mal_predicted_diff)

attack_name_list.append(attack)
eps_val_list.append(eps_val)

```

0%| | 0/5 [00:00<?, ?it/s]

```

[12]: df_result = pd.DataFrame({'Malicious_Distance':mal_diff_distance_list,
                                'Real_Predicted_MSE':real_mse_list,
                                'Malicious_Predicted_MSE':mal_mse_list,
                                'MalOut_RealOut_Diff':mal_predicted_diff_list,
                                'Attack':attack_name_list,
                                'eps':eps_val_list})

df_result

```

```

[12]:
   Malicious_Distance  Real_Predicted_MSE  Malicious_Predicted_MSE  \
0          281.642853          0.007904          0.010725
1           81.320984          0.007904          0.008754
2           85.205322          0.007904          0.008792
3           14.572388          0.007904          0.007941
4          196.103790          0.007904          0.010890
..          ...
115         198.043716          0.006840          0.008432
116          34.424145          0.006840          0.006948
117         207.769989          0.006840          0.007591
118         268.659790          0.006840          0.007416
119         236.913177          0.006840          0.007804

MalOut_RealOut_Diff  Attack  eps

```


|     |           |      |     |
|-----|-----------|------|-----|
| 0   | 20.035965 | MIM  | 0.8 |
| 1   | 9.123885  | PGD  | 0.1 |
| 2   | 8.614745  | MIM  | 0.1 |
| 3   | 1.382549  | FGSM | 0.3 |
| 4   | 17.165459 | PGD  | 0.5 |
| ..  | ...       | ...  | ... |
| 115 | 18.537216 | PGD  | 0.3 |
| 116 | 3.128234  | FGSM | 0.7 |
| 117 | 15.462641 | BIM  | 0.3 |
| 118 | 20.827293 | BIM  | 0.8 |
| 119 | 17.289284 | PGD  | 0.7 |

[120 rows x 6 columns]

```
[13]: fig, ax = plt.subplots(1,2, sharey=True, figsize=(15,5))
sns.lineplot(x='eps', y='Real_Predicted_MSE', hue='Attack', ci='sd',
             estimator="median",
             data=df_result, ax=ax[0])

sns.lineplot(x='eps', y='Malicious_Predicted_MSE', hue='Attack', ci='sd',
             estimator="median",
             data=df_result, ax=ax[1])
```

[13]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f94a060ba10>



#Defensive Distillation

```
[14]: !rm -f util_defdistill.py
!wget -q https://raw.githubusercontent.com/ocatak/6g_MIMO_security_distillation/
      main/util_defdistill.py

from util_defdistill import Distiller
```

```

STUDENT_MODEL_MULTIPLICATION = 0.5
TEACHER_MODEL_MULTIPLICATION = 1.0

student_model = get_model(STUDENT_MODEL_MULTIPLICATION)
teacher_model = get_model(TEACHER_MODEL_MULTIPLICATION)

es_teacher = EarlyStopping(monitor='val_loss',
                           patience=150,
                           min_delta=0.001,
                           verbose=1,
                           restore_best_weights=True,
                           mode='min')

# Train and evaluate teacher on data.
hist_teacher = teacher_model.fit(X_train, y_train,
                                epochs=100000,
                                verbose=0, batch_size=1024,
                                callbacks=[es_teacher, TqdmCallback(verbose=1, desc='DL_
↪training')],
                                validation_data=(X_test, y_test))

```

DL training: 0epoch [00:00, ?epoch/s]

DL training: 0batch [00:00, ?batch/s]

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0245s vs `on\_train\_batch\_end` time: 0.0301s). Check your callbacks.

Restoring model weights from the end of the best epoch: 405.


Epoch 555: early stopping

```

[15]: plot_history(hist_teacher.history)
      plt.show()

```





```
[16]: import keras
distiller = Distiller(student=student_model, teacher=teacher_model)
loss_fn = keras.losses.MeanSquaredError()

distiller.compile(optimizer='adam',
                  metrics=['mse'],
                  student_loss_fn=loss_fn,
                  distillation_loss_fn=keras.losses.KLDivergence(),
                  alpha=0.1,
                  temperature=20)

es = EarlyStopping(monitor='val_student_loss',
                  patience=20,
                  #min_delta=0.000001,
                  verbose=1,
                  restore_best_weights=True,
                  mode='min')

hist_distill = distiller.fit(X_train, y_train,
                             epochs=10000,
                             verbose=0,
                             callbacks=[es, TqdmCallback(verbose=1, desc='DL_
→training')],
                             validation_data=(X_test, y_test))
```


DL training: 0epoch [00:00, ?epoch/s]

DL training: 0batch [00:00, ?batch/s]

Restoring model weights from the end of the best epoch: 37.

Epoch 57: early stopping

```
[17]: plot_history(hist_distill.history)
      plt.show()
```



```
[18]: from keras.layers import Dense, Conv2D, Flatten
      student_copy = tf.keras.models.clone_model(distiller.student)
      for layer in student_copy.layers:
          layer.trainable = False


      student_copy.add(Dense(2, activation='softmax'))
      student_copy.compile(optimizer='adam', loss='categorical_crossentropy',
                           metrics=['accuracy'])

      y_dummy = np.zeros((X_train.shape[0],2))
      mid_range = np.int(X_train.shape[0] / 2.0)
      y_dummy[0:mid_range,0]= 1.0
      history = student_copy.fit(X_train, y_dummy, epochs=2, verbose=0,
                                callbacks=[TqdmCallback(verbose=1)],batch_size=128,
                                validation_split=0.33, shuffle=True)
      plot_history(history)
      plt.show()
```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:10:  
DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To  
silence this warning, use `int` by itself. Doing this will not modify any  
behavior and is safe. When replacing `np.int`, you may wish to use e.g.  
`np.int64` or `np.int32` to specify the precision. If you wish to review your  
current use, check the release note link for additional information.  
Deprecated in NumPy 1.20; for more details and guidance:  
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>  
# Remove the CWD from sys.path while we load stuff.

0epoch [00:00, ?epoch/s]

0batch [00:00, ?batch/s]



```
[19]: from itertools import product
      from random import shuffle

      logits_model_student = tf.keras.Model(student_copy.input, student_copy.output)
      attacks = ['FGSM', 'BIM', 'MIM', 'PGD', 'CW']
      attacks = ['FGSM', 'BIM', 'MIM', 'PGD']

      eps_vals = [0.01, 0.1, 0.3, 0.5, 0.7, 0.8]

      params_list = list(product(eps_vals, attacks))
      shuffle(params_list)

      distil_mal_diff_distance_list = []
      distil_real_mse_list = []
      distil_mal_mse_list = []
      distil_mal_predicted_diff_list = []
      distil_attack_name_list = []
      distil_eps_val_list = []

      #for i in tqdm(range(X_test.shape[0])):
      for _ in tqdm(range(5)):
          i = np.random.randint(0, X_test.shape[0])
          test_input = X_test[i:i+1, :].astype(np.float32)

          for eps_val, attack in params_list:
              if attack == 'FGSM':
```

```

        mal_input = fast_gradient_method(logits_model, test_input,
↪eps=eps_val/20.0, norm=np.inf,
                                clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                targeted=False)

        mal_input = mal_input.numpy()
        elif attack == 'BIM':
            mal_input = basic_iterative_method(logits_model, test_input,
↪eps=eps_val, norm=np.inf,
                                #clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                eps_iter=0.01,nb_iter=2000,
                                targeted=False)

            mal_input = mal_input.numpy()
            elif attack == 'MIM':
                mal_input = momentum_iterative_method(logits_model, test_input,
↪eps=eps_val, norm=np.inf,
                                clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                eps_iter=0.01,nb_iter=2000,
                                targeted=False)

                mal_input = mal_input.numpy()
                elif attack == 'PGD':
                    mal_input = projected_gradient_descent(logits_model, test_input,
↪eps=eps_val, norm=np.inf,
                                clip_min=test_input.
↪min(),clip_max=test_input.max(),
                                eps_iter=0.01,nb_iter=2000,
                                targeted=False)

                    mal_input = mal_input.numpy()
                    elif attack == 'CW':
                        mal_input = carlini_wagner_l2(logits_model, test_input,
                                targeted=False, #y=y_output,
                                batch_size=512, confidence=10.0,
                                abort_early=False, max_iterations=1000,
                                clip_min=test_input.min(),clip_max=test_input.
↪max())

        mal_diff = mal_input - test_input
        mal_diff_distance = norm(mal_diff, ord = np.inf)
        distil_mal_diff_distance_list.append(mal_diff_distance)

        real_output = y_test[i:i+1,:]
        test_output = distiller.student.predict(test_input)
        real_mse = mean_squared_error(real_output,test_output)
        distil_real_mse_list.append(real_mse)

```

```

malicious_test_output = distiller.student.predict(mal_input)
mal_mse = mean_squared_error(real_output, malicious_test_output)
distil_mal_mse_list.append(mal_mse)

mal_predicted_diff = malicious_test_output - test_output
mal_predicted_diff = norm(mal_predicted_diff, ord = np.inf)
distil_mal_predicted_diff_list.append(mal_predicted_diff)

distil_attack_name_list.append(attack)
distil_eps_val_list.append(eps_val)

```

0%| | 0/5 [00:00<?, ?it/s]

```

[20]: df_result_distill = pd.DataFrame({'Malicious_Distance':
    ↪ distil_mal_diff_distance_list,
                                     'Real_Predicted_MSE':distil_real_mse_list,
                                     'Malicious_Predicted_MSE':distil_mal_mse_list,
                                     'MalOut_RealOut_Diff':distil_mal_predicted_diff_list,
                                     'Attack':distil_attack_name_list,
                                     'eps':distil_eps_val_list})
df_result_distill

```

```

[20]:
   Malicious_Distance  Real_Predicted MSE  Malicious_Predicted MSE  \
0          90.840012          0.000406          0.000408
1         205.466568          0.000406          0.000596
2           9.849995          0.000406          0.000404
3         334.129913          0.000406          0.000455
4           5.057673          0.000406          0.000408
..          ...          ...          ...
115         197.000610          0.000243          0.008216
116         192.487305          0.000243          0.007124
117         259.968994          0.000243          0.008748
118           9.880083          0.000243          0.000237
119         200.001724          0.000243          0.009183

   MalOut_RealOut_Diff  Attack  eps
0          3.866464      BIM  0.10
1          8.485148      PGD  0.30
2          0.130710      BIM  0.01
3          9.804306      BIM  0.70
4          0.110424     FGSM  0.10
..          ...          ...    ...
115         38.448257      PGD  0.80
116         34.984123      MIM  0.30
117         40.053532      MIM  0.50
118          0.343055      MIM  0.01

```

119                      41.169380      PGD   0.50

[120 rows x 6 columns]

```
[21]: fig, ax = plt.subplots(1,2, sharey=True, figsize=(15,5))
sns.lineplot(x='eps', y='Malicious_Predicted_MSE', hue='Attack', ci='sd',
             ↪estimator="median",
             data=df_result, ax=ax[0])

sns.lineplot(x='eps', y='Malicious_Predicted_MSE', hue='Attack', ci='sd',
             ↪estimator="median",
             data=df_result_distill, ax=ax[1])
```

[21]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f94a0350a90>

